

Skygeni Assignment

Submitted by Jeevana Sree B (12220183)

I have been given 4 csv files -

subscription_information.csv, payment_information.csv, industry_clients_details.csv and financial_information.csv. This assignment involves analyzing this four datasets to derive meaningful insights and find answers to the given questions.

Importing necessary libraries

```
import pandas as pd # data handling and manipulation
import matplotlib.pyplot as plt # Basic Visualization
import seaborn as sns # Advanced Visualization
```

Loading Data

```
financial_df = pd.read_csv('financial_information.csv')
industry_df = pd.read_csv('industry_client_details.csv')
payments_df = pd.read_csv('payment_information.csv')
subscription_df = pd.read_csv('subscription_information.csv')
```

Data Inspection and Cleaning

1. Financial Data

```
# display the first 5 rows
financial_df.head()
```

	Unnamed: 0	start_date	end_date	inflation_rate	gdp_growth_rate
0	0	2018-01-01	2018-03-31	5.77	3.51
1	1	2018-04-01	2018-06-30	1.17	2.15
2	2	2018-07-01	2018-09-30	1.56	1.82
3	3	2018-10-01	2018-12-31	2.78	2.43
4	4	2019-01-01	2019-03-31	6.91	3.44

```
# No of rows and columns
financial_df.shape
```

```
(21, 5)
```

```
# information about dataframe
financial_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21 entries, 0 to 20
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      21 non-null    int64
```

```

1  start_date      21 non-null    object
2  end_date        21 non-null    object
3  inflation_rate  21 non-null    float64
4  gdp_growth_rate 21 non-null    float64
dtypes: float64(2), int64(1), object(2)
memory usage: 968.0+ bytes

# to check columns
financial_df.columns.tolist()

['Unnamed: 0', 'start_date', 'end_date', 'inflation_rate',
'gdp_growth_rate']

# Statistical summary
financial_df.describe()

```

	Unnamed: 0	inflation_rate	gdp_growth_rate
count	21.000000	21.000000	21.000000
mean	10.000000	4.426667	2.345238
std	6.204837	2.245801	0.952747
min	0.000000	0.760000	1.040000
25%	5.000000	2.710000	1.360000
50%	10.000000	4.400000	2.400000
75%	15.000000	6.760000	3.440000
max	20.000000	7.710000	3.630000

- the financial data contains 21 rows and 5 columns
- `start_date` and `end_date` says about the financial periods (in string data type)
- `inflation_rate` and `gdp_growth_rate` represents the inflation and gdp growth during the financial periods
- `unnamed` it looks like the index column of the data frame

```

# Check for any missing values
financial_df.isnull().sum()

Unnamed: 0      0
start_date      0
end_date        0
inflation_rate  0
gdp_growth_rate 0
dtype: int64

# check for duplicate values
duplicate_count = financial_df.duplicated().sum()
duplicate_count

0

```

There are no missing values and duplicate values. So, we can move on the handle `unnamed: 0` column and converting `start_date`, `end_date` columns into datetime format

```
# Since unnamed column also says about the index of the data frame it
better to remove the unnecessary column
# removing 'unnamed: 0' column
financial_df.drop(columns=['Unnamed: 0'], inplace=True)
```

inplace = True modifies the original dataframe directly instead of creating a new one.

```
# Convert 'start_date' and 'end_date' columns to datetime format
financial_df['start_date'] =
pd.to_datetime(financial_df['start_date'])
financial_df['end_date'] = pd.to_datetime(financial_df['end_date'])

# Ensure 'inflation_rate' and 'gdp_growth_rate' are clean floats
column
financial_df['inflation_rate'] =
pd.to_numeric(financial_df['inflation_rate'], errors='coerce')
financial_df['gdp_growth_rate'] =
pd.to_numeric(financial_df['gdp_growth_rate'], errors='coerce')

# Show the cleaned DataFrame info and first few rows
financial_df.info(), financial_df.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21 entries, 0 to 20
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   start_date            21 non-null    datetime64[ns]
1   end_date              21 non-null    datetime64[ns]
2   inflation_rate        21 non-null    float64
3   gdp_growth_rate       21 non-null    float64
dtypes: datetime64[ns](2), float64(2)
memory usage: 800.0 bytes
```

```
(None,
  start_date  end_date  inflation_rate  gdp_growth_rate
0 2018-01-01 2018-03-31          5.77          3.51
1 2018-04-01 2018-06-30          1.17          2.15
2 2018-07-01 2018-09-30          1.56          1.82
3 2018-10-01 2018-12-31          2.78          2.43
4 2019-01-01 2019-03-31          6.91          3.44)
```

2. Industry Clienta Dataset

```
# show first 5 rows()
industry_df.head()
```

	client_id	company_size	industry	location
0	4280387012	Large	Finance Lending	Mumbai
1	2095513148	Small	Finance Lending	Chennai

2	7225516707	Medium	Finance Lending	New Delhi
3	8093537819	Large	Block Chain	Mumbai
4	4387541014	Medium	Hyper Local	Bangalore

```
# information about dataset
```

```
industry_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 100 entries, 0 to 99
```

```
Data columns (total 4 columns):
```

#	Column	Non-Null Count	Dtype
0	client_id	100 non-null	int64
1	company_size	100 non-null	object
2	industry	100 non-null	object
3	location	100 non-null	object

```
dtypes: int64(1), object(3)
```

```
memory usage: 3.2+ KB
```

```
# No of rows and Columns
```

```
industry_df.shape
```

```
(100, 4)
```

The dataset contains 100 rows and 4 columns

```
# Checking for missing values
```

```
industry_df.isnull().sum()
```

```
client_id      0
```

```
company_size   0
```

```
industry        0
```

```
location        0
```

```
dtype: int64
```

```
# checking for duplicates
```

```
industry_df.duplicated().sum()
```

```
0
```

```
# Ensures consistency across datasets during merging, even if already int64
```

```
industry_df['client_id'] = industry_df['client_id'].astype('int64')
```

```
# Cleaning categorical columns to remove whitespace and standardize format using title case
```

```
industry_df['industry'] =
```

```
industry_df['industry'].str.strip().str.title()
```

```
industry_df['company_size'] =
```

```
industry_df['company_size'].str.strip().str.title()
```

```
industry_df['location'] =  
industry_df['location'].str.strip().str.title()
```

```
# Display Cleaned Dataset
```

```
industry_df.head()
```

	client_id	company_size	industry	location
0	4280387012	Large	Finance Lending	Mumbai
1	2095513148	Small	Finance Lending	Chennai
2	7225516707	Medium	Finance Lending	New Delhi
3	8093537819	Large	Block Chain	Mumbai
4	4387541014	Medium	Hyper Local	Banglore

```
# Categorical Value Distribution
```

```
# Company size distribution
```

```
print("\nCompany Size Distribution:")
```

```
print(industry_df['company_size'].value_counts())
```

```
# Industry distribution
```

```
print("\nIndustry Distribution:")
```

```
print(industry_df['industry'].value_counts())
```

```
# Location distribution
```

```
print("\nLocation Distribution:")
```

```
print(industry_df['location'].value_counts())
```

```
Company Size Distribution:
```

```
Small      39
```

```
Medium     32
```

```
Large      29
```

```
Name: company_size, dtype: int64
```

```
Industry Distribution:
```

```
Block Chain      25
```

```
Finance Lending  22
```

```
Gaming           22
```

```
Hyper Local      20
```

```
Ai               11
```

```
Name: industry, dtype: int64
```

```
Location Distribution:
```

```
Hyderabad      23
```

```
Chennai        20
```

```
New Delhi      20
```

```
Banglore       20
```

```
Mumbai         17
```

```
Name: location, dtype: int64
```

Information about the columns

1. `client_id` - it represents unique value assigned to each client (PRIMARY KEY)
2. `company_size` - indicates the scale of company (small, medium and large scales)
3. `industry` - information about the sector or industry the client operates
4. `location` - Says about the city the client is located

The industry dataset have no missing values, no duplicates. I have cleaned categorical columns to remove whitespace and standardize format using title case.

How many finance lending and blockchain clients does the organization have? The organisation have 25 Block Chain Clients and 22 Finance Lending Clients. (Check in the Value counts)

3. Subscription Information Dataset

```
# display first 5 rows
subscription_df.head()
```

	client_id	subscription_type	start_date	end_date	renewed
0	1131383004	Yearly	2020-11-11	2021-11-11	False
1	4309371709	Monthly	2021-05-24	2021-06-23	True
2	3183675157	Yearly	2021-12-25	2022-12-25	True
3	5371694837	Monthly	2020-03-14	2020-04-13	True
4	5157113076	Monthly	2019-11-07	2019-12-07	False

```
# no of rows and Columns
subscription_df.shape
```

```
(100, 5)
```

```
# information about dataset
subscription_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   client_id              100 non-null   int64
1   subscription_type       100 non-null   object
2   start_date             100 non-null   object
3   end_date               100 non-null   object
4   renewed                100 non-null   bool
dtypes: bool(1), int64(1), object(3)
memory usage: 3.3+ KB
```

```
# Checking for missing values
subscription_df.isnull().sum()
```

```
client_id      0
subscription_type  0
start_date     0
end_date       0
```

```

renewed          0
dtype: int64

# checking for duplicate values
subscription_df.duplicated().sum()

0

# Ensure client_id is of type int64
subscription_df['client_id'] =
subscription_df['client_id'].astype('int64')

# Convert start_date and end_date columns to datetime format
subscription_df['start_date'] =
pd.to_datetime(subscription_df['start_date'])
subscription_df['end_date'] =
pd.to_datetime(subscription_df['end_date'])

# Verify data types after cleaning
print("\nData Types After Cleaning:")
print(subscription_df.dtypes)

Data Types After Cleaning:
client_id          int64
subscription_type  object
start_date        datetime64[ns]
end_date          datetime64[ns]
renewed            bool
dtype: object

# Subscription Types distribution
print("\nSubscription type count:")
print(subscription_df['subscription_type'].value_counts())

# Renewed or Not Distribution
print("\nRenewal Count:")
print(subscription_df['renewed'].value_counts())

Subscription type count:
Monthly    57
Yearly     43
Name: subscription_type, dtype: int64

Renewal Count:
True       55
False      45
Name: renewed, dtype: int64

# Display cleaned dataset
subscription_df.head()

```

	client_id	subscription_type	start_date	end_date	renewed
0	1131383004	Yearly	2020-11-11	2021-11-11	False
1	4309371709	Monthly	2021-05-24	2021-06-23	True
2	3183675157	Yearly	2021-12-25	2022-12-25	True
3	5371694837	Monthly	2020-03-14	2020-04-13	True
4	5157113076	Monthly	2019-11-07	2019-12-07	False

Summary

1. `client_id` - Unique ID for each client (used to link with other datasets)
2. `subscription_type` - Type of subscription (Monthly or Yearly)
3. `start_date` - Date when the subscription started
4. `end_date` - Date when the subscription ended
5. `renewed` - True if the client renewed the subscription, else False

The subscription dataset has no missing and duplicate values. And converted start_date, end_date columns into datetime format

4. Payments Dataset

```
# show first 5 rows
payments_df.head()
```

	client_id	payment_date	amount_paid	payment_method
0	6292156167	9/16/2019	447.0	Bank Transfer
1	7462725203	5/21/2018	379.7	Bank Transfer
2	4698004907	9/11/2021	435.1	Check
3	3510240337	12/7/2020	413.1	Check
4	7501599785	3/4/2019	61.1	Bank Transfer

```
# Information about dataset
payments_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   client_id             100 non-null    int64
1   payment_date          100 non-null    object
2   amount_paid           100 non-null    float64
3   payment_method        100 non-null    object
dtypes: float64(1), int64(1), object(2)
memory usage: 3.2+ KB
```

```
# checking shape
payments_df.shape
```

```
(100, 4)
```



```

# Checking for null values
payments_df.isnull().sum()

client_id      0
payment_date   0
amount_paid    0
payment_method 0
dtype: int64

# Checking for duplicate values
payments_df.duplicated().sum()

0

# Ensure client_id is of type int64
payments_df['client_id'] = payments_df['client_id'].astype('int64')

# Convert payment_date to datetime format
payments_df['payment_date'] =
pd.to_datetime(payments_df['payment_date'])

# Clean and standardize payment_method (strip spaces, apply title
case)
payments_df['payment_method'] =
payments_df['payment_method'].str.strip().str.title()

# Ensure amount_paid is a clean float column
payment_df['amount_paid'] = pd.to_numeric(payment_df['amount_paid'],
errors='coerce')

# Verify data types after cleaning
payment_df.dtypes

client_id      int64
payment_date   object
amount_paid    float64
payment_method  object
dtype: object

# Payments Method Distribution
print("\nTypes of payment methods:")
print(payments_df['payment_method'].value_counts())

Types of payment methods:
Bank Transfer    38
Check            31
Credit Card     31
Name: payment_method, dtype: int64

# Display Cleaned Data Set
payment_df.head()

```

	client_id	payment_date	amount_paid	payment_method
0	6292156167	9/16/2019	447.0	Bank Transfer
1	7462725203	5/21/2018	379.7	Bank Transfer
2	4698004907	9/11/2021	435.1	Check
3	3510240337	12/7/2020	413.1	Check
4	7501599785	3/4/2019	61.1	Bank Transfer

Now we are done with Data Loading, Cleaning and Transforming, the datasets can be used for Data Visualization.

With the help of the plots, We can answer the questions.

Question 1: How many finance lending and blockchain clients does the organization have?

Approach:

1. Using industry_df
2. Filter rows where industry is "Finance Lending" or "Block Chain"
3. Count them
4. Visualize using a bar plot for better clarity

```
# Filter the dataset for required industries
target_industries = ['Finance Lending', 'Block Chain']
filtered_clients =
industry_df[industry_df['industry'].isin(target_industries)]


# Count number of clients in each industry
industry_counts = filtered_clients['industry'].value_counts()

# Plotting the result
plt.figure(figsize=(6, 4))
ax = sns.barplot(x=industry_counts.index, y=industry_counts.values)

# Add renewal rate labels on top of bars
for i, v in enumerate(industry_counts.values):
    ax.text(i, v+0.2, str(v), ha='center', fontsize=10)

plt.title("Number of Clients in Finance Lending & Blockchain Industries")
plt.xlabel("Industry")
plt.ylabel("Number of Clients")
plt.show()

# Print count values for report
total_clients = industry_counts.sum()
print("Total Target Clients (Finance + Blockchain):", total_clients)
print("\nBreakdown:")
print(industry_counts)
```



Total Target Clients (Finance + Blockchain): 47

Breakdown:

Block Chain 25

Finance Lending 22

Name: industry, dtype: int64

Question 2: Which industry in the organization has the highest renewal rate?

Approach:

1. Will Merge datasets: subscription_df + industry_df on `clients_id`
2. Group by industry and calculate: Renewal Rate =
(Number of Renewed Clients/Total Clients in that Industry)
3. Sort industries by renewal rate
4. Visualize using a bar plot with value labels

```
# Merge subscription and industry data on client_id
merged = pd.merge(subscription_df, industry_df, on='client_id')

# Group by industry and calculate average of 'renewed' (which gives
the renewal rate)
renewal_rate = merged.groupby('industry')
['renewed'].mean().sort_values(ascending=False)

# Plot using seaborn with value labels
plt.figure(figsize=(8, 5))
```

```

ax = sns.barplot(x=renewal_rate.index, y=renewal_rate.values)

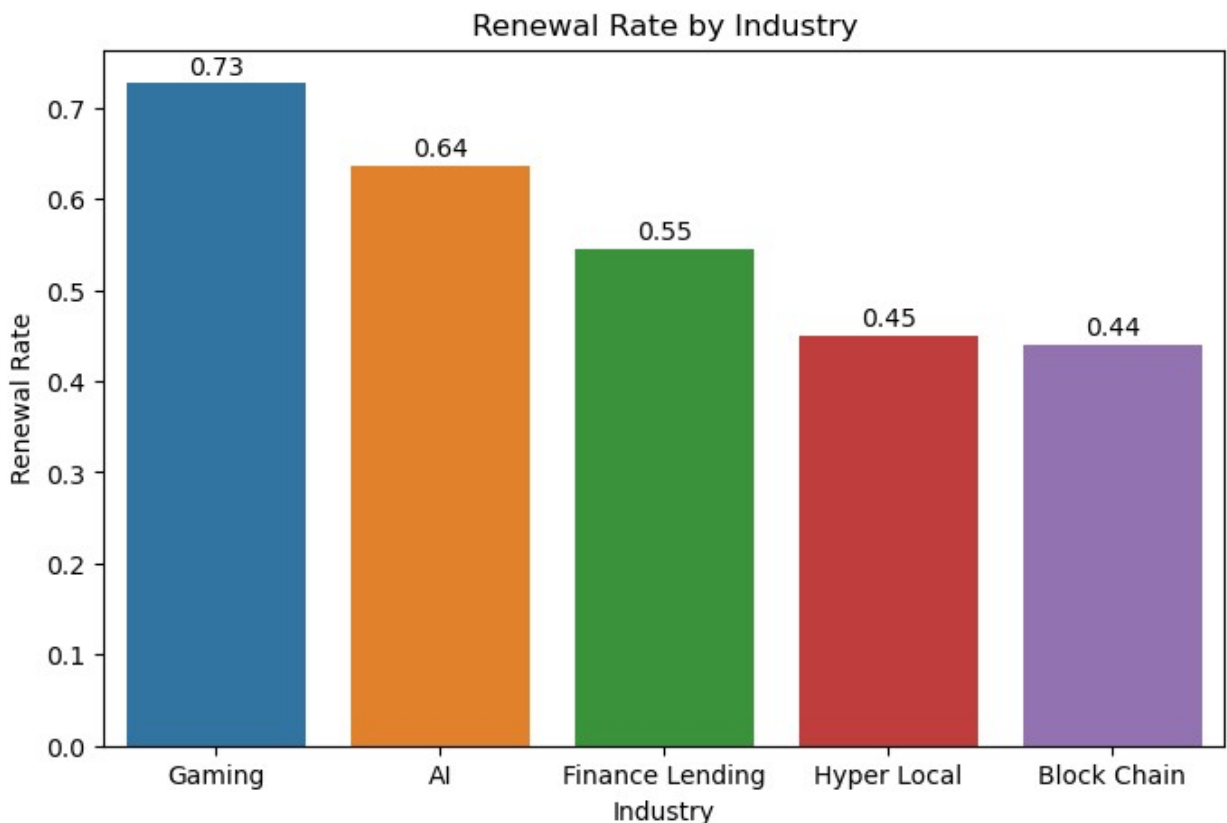
# Add renewal rate labels on top of bars
for i, v in enumerate(renewal_rate.values):
    ax.text(i, v + 0.01, f"{v:.2f}", ha='center', fontsize=10)

plt.title("Renewal Rate by Industry")
plt.xlabel("Industry")
plt.ylabel("Renewal Rate")
plt.show()

# Display the industry with the highest renewal rate
top_industry = renewal_rate.idxmax()
top_rate = renewal_rate.max()

print(f"The industry with the highest renewal rate is {top_industry} with {top_rate:.2f}")

```



The industry with the highest renewal rate is Gaming with 0.73

Question 3: What was the average inflation rate when their subscriptions were renewed?

Approach:

1. From `subscription_df`, filter rows where `renewed == True`
2. For each renewed subscription, find the matching row in `financial_df` where: The subscription's `start_date` (or `end_date`) falls between the financial `start_date` and `end_date`
3. From those matched financial periods, get the `inflation_rate`
4. Calculate the average inflation rate for all renewed subscriptions

```
# Keep only the subscriptions that were renewed
renewed_data = subscription_df[subscription_df['renewed'] ==
True].copy()

# Add a new column to store matching inflation rates
renewed_data['inflation_rate'] = None

# Go through each renewed subscription and find its matching financial
period
for i in renewed_data.index:
    sub_date = renewed_data.loc[i, 'start_date'] # get the
subscription start date

    # Find where this date fits in the financial period
    match = financial_df[
        (financial_df['start_date'] <= sub_date) &
        (financial_df['end_date'] >= sub_date)
    ]

    # If we find a match, save the inflation rate
    if not match.empty:
        renewed_data.at[i, 'inflation_rate'] = match.iloc[0]
['inflation_rate']

# Make sure inflation_rate is treated as a number
renewed_data['inflation_rate'] =
pd.to_numeric(renewed_data['inflation_rate'])

# Calculate the average inflation rate
average_inflation = renewed_data['inflation_rate'].mean()

# Print the result
print("Average inflation rate during renewals:",
round(average_inflation, 2), "%")

# Show a few sample rows
print(renewed_data[['client_id', 'start_date',
'inflation_rate']].head())
```

```
Average inflation rate during renewals: 4.44 %
   client_id  start_date  inflation_rate
1  4309371709  2021-05-24             0.76
2  3183675157  2021-12-25             7.32
```

3	5371694837	2020-03-14	4.40
5	7896208406	2022-02-24	6.76
6	4687291312	2019-06-14	3.84

- `.copy()`: Makes a separate copy of the data so changes don't affect the original.
- `.loc[]`: Lets us get or update values using labels (like column names).
- `.iloc[]`: Lets us get values using row numbers (positions).

Question 4: What is the median amount paid each year for all payment methods?

Approach:

1. Use the `payments_df` dataset
2. Extract year from the `payment_date` column
3. Group by `year` and `payment_method`
4. Calculate the median of `amount_paid` for each group
5. Visualize using a grouped bar plot by year & payment method

```
# Add a new column to extract the year from payments_date
payments_df['year'] = payments_df['payment_date'].dt.year

# Group by year and payments method, then find the median amount paid
median_by_year = payments_df.groupby(['year', 'payment_method'])
['amount_paid'].median().reset_index()

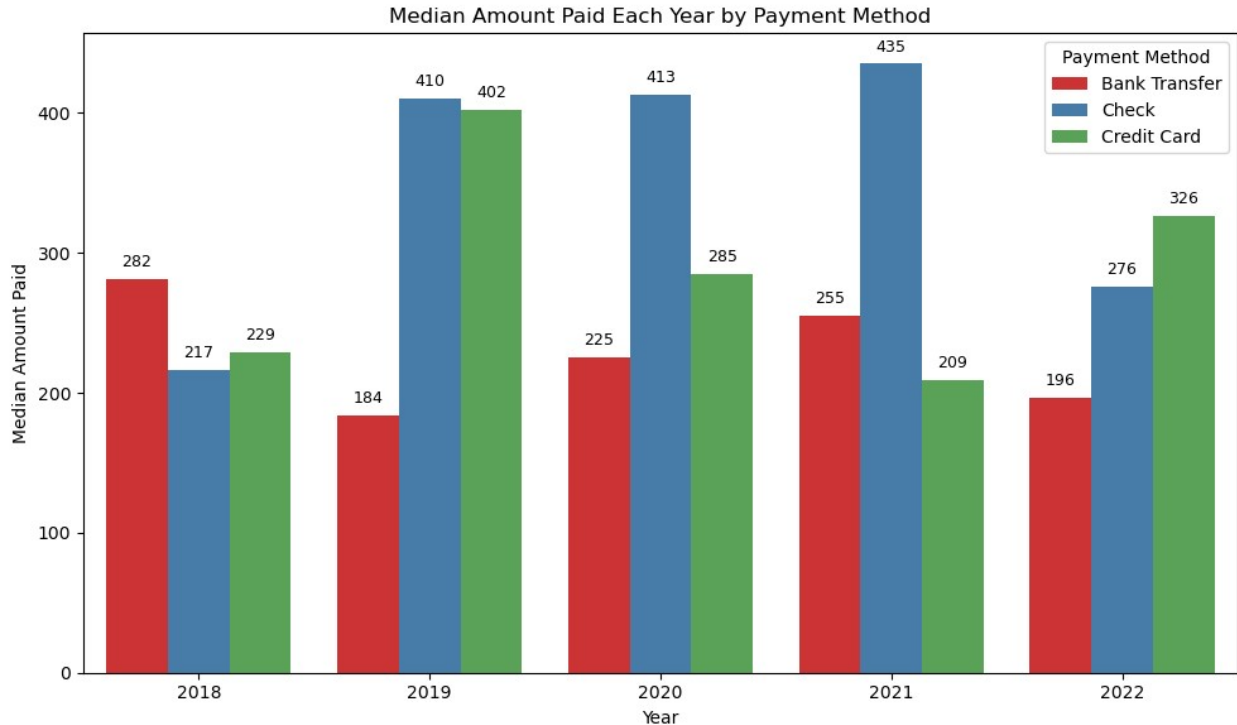
# Set the size of the chart
plt.figure(figsize=(10, 6))

# Create the barplot
ax = sns.barplot(
    data=median_by_year,
    x='year',
    y='amount_paid',
    hue='payment_method',
    palette='Set1'
)

# Add value labels on top of each bar
for bar in ax.patches:
    height = bar.get_height()
    ax.annotate(f'{height:.0f}', # value rounded to 0 decimals
                xy=(bar.get_x() + bar.get_width() / 2, height), #
                label position
                xytext=(0, 5), # offset
                textcoords='offset points',
                ha='center', va='bottom', fontsize=9)

# Set labels and title
plt.title("Median Amount Paid Each Year by Payment Method")
plt.xlabel("Year")
```

```
plt.ylabel("Median Amount Paid")
plt.legend(title="Payment Method")
plt.tight_layout()
plt.show()
```



Adding a separate `year` column makes it easy to group and compare payments by year. It also makes the chart cleaner and the x-axis labels clearer.

Below are the median payment amounts made by clients for each payment method, year by year:

Year	Payment Method	Median Amount Paid
2018	Bank Transfer	282
2018	Check	217
2018	Credit Card	229
2019	Bank Transfer	184
2019	Check	410
2019	Credit Card	402
2020	Bank Transfer	225
2020	Check	413
2020	Credit Card	285
2021	Bank Transfer	255
2021	Check	435

Year	Payment Method	Median Amount Paid
2021	Credit Card	209
2022	Bank Transfer	196
2022	Check	276
2022	Credit Card	326